

# Sim-to-Real

## An Incomplete Overview


Jie Tan

CS 8803 Deep Reinforcement Learning for Intelligent Control

04/04/2022

# What's the sim-to-real gap?

Dynamics:



Perception:



# Goals

- Understand the causes of sim-to-real
- Review of the state-of-the-art methods
  - System Identification
  - Domain Randomization
  - Domain Adaptation
  - Meta Learning

# Brainstorming: Why Sim-to-Real?

## Real world

- Slow
- Unsafe
- Expensive
- Human supervision


## Simulation

- Fast
- Safe
- Cheap
- Scalable

# Why Sim-to-Real?

Transform hard robotic problems into large-scale computation problems

# Trend on Sim-to-Real



# How to overcome sim-to-real gap?

- Improve simulation
  - System identification
    - [Sim-to-Real: Learning Agile Locomotion For Quadruped Robots](#)
    - [Simulation-Based Design of Dynamic Controllers for Humanoid Balancing](#)
    - [Preparing for the Unknown: Learning a Universal Policy with Online System Identification](#)
- Improve policy
  - Domain randomization
    - [Sim-to-Real Transfer of Robotic Control with Dynamics Randomization](#)
    - [Closing the Sim-to-Real Loop: Adapting Simulation Randomization with Real World Experience](#)
  - Domain adaptation
    - [Learning Agile Robotic Locomotion Skills by Imitating Animals](#)
    - [Sim-to-Real Transfer for Biped Locomotion](#)
  - Meta learning
    - [Rapidly Adaptable Legged Robots via Evolutionary Meta-Learning](#)
    - [NoRML: No Reward Meta-Learning](#)
    - [Policy Transfer with Strategy Optimization](#)

# System Identification




# What are the causes of sim-to-real gap?

- Unmodeled dynamics
- Wrong simulation parameters
- Inaccurate contact models
- Latency
- Actuator dynamics
- Noise
- Stochastic real environment
- Numerical accuracy
- ...




**Actuator dynamics** and **latency** are two important causes of reality gap.



[\[Sim-to-Real: Learning Agile Locomotion For Quadruped Robots, RSS 2018\]](#)

# Actuator Model

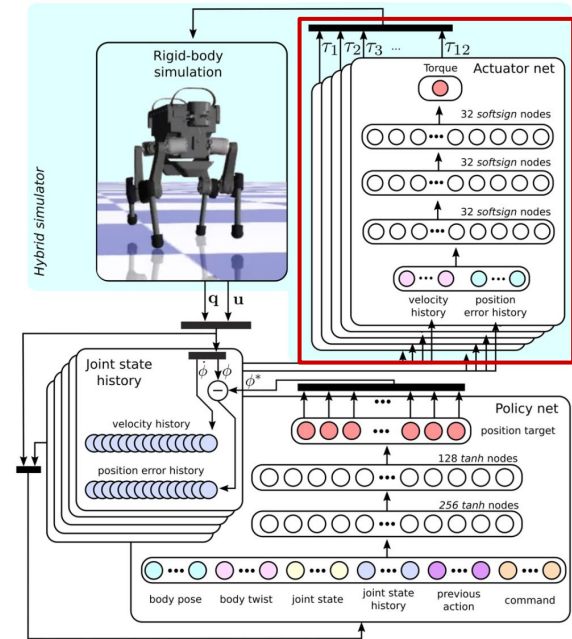
## Analytical models




$$\begin{aligned}\tau &= f(I) \\ I &= \frac{V * \text{PWM} - V_{\text{emf}}}{R} \\ V_{\text{emf}} &= K_t \dot{q}\end{aligned}$$

[\[Sim-to-Real: Learning Agile Locomotion For Quadruped Robots, RSS 2018\]](#)

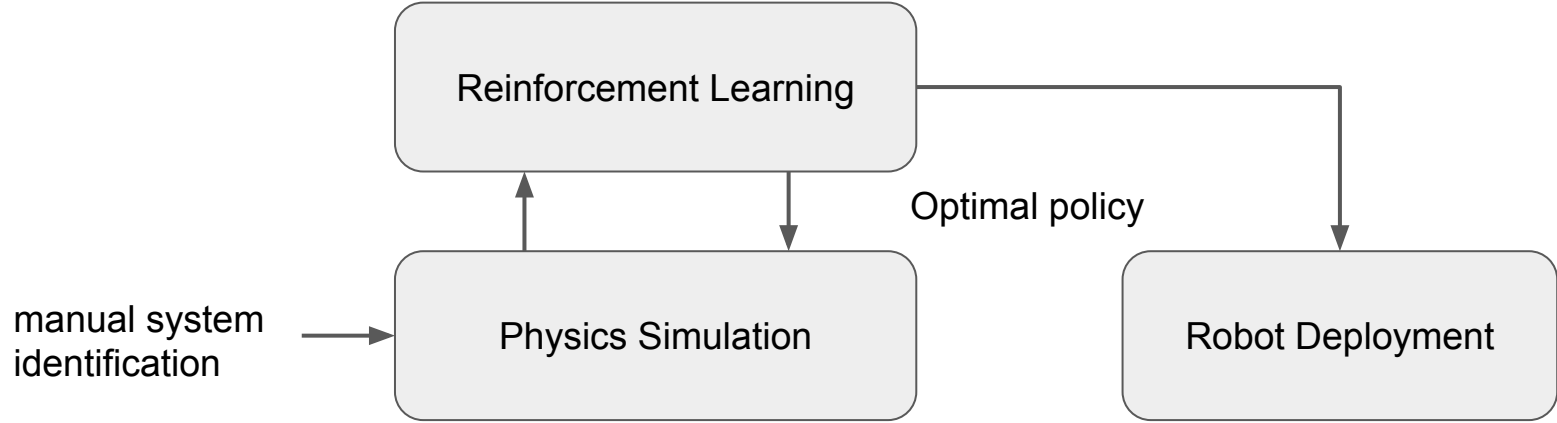
## Neural network models



[\[Learning agile and dynamic motor skills for legged robots, Science Robotics 2019\]](#)





[\[Sim-to-Real: Learning Agile Locomotion For Quadruped Robots, RSS 2018\]](#)

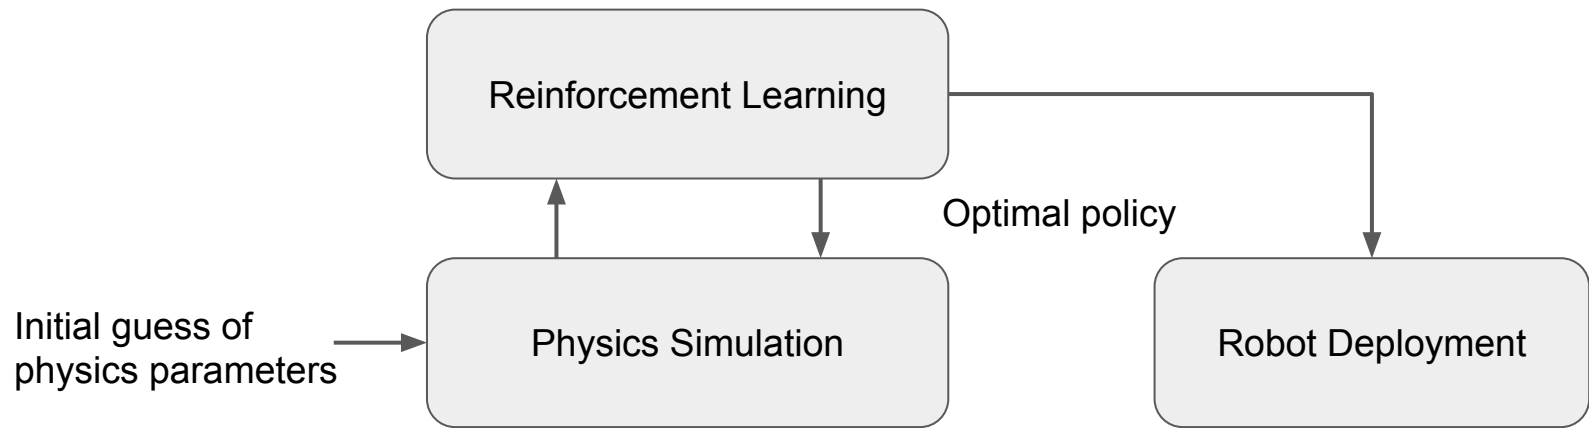


- Limitations

- Disassemble the robot
- Decide what parameters to identify
- Design experiments for individual parameters
- Lots of manual work



[\[Simulation-based design of dynamic controllers for humanoid balancing, IROS 2016\]](#)

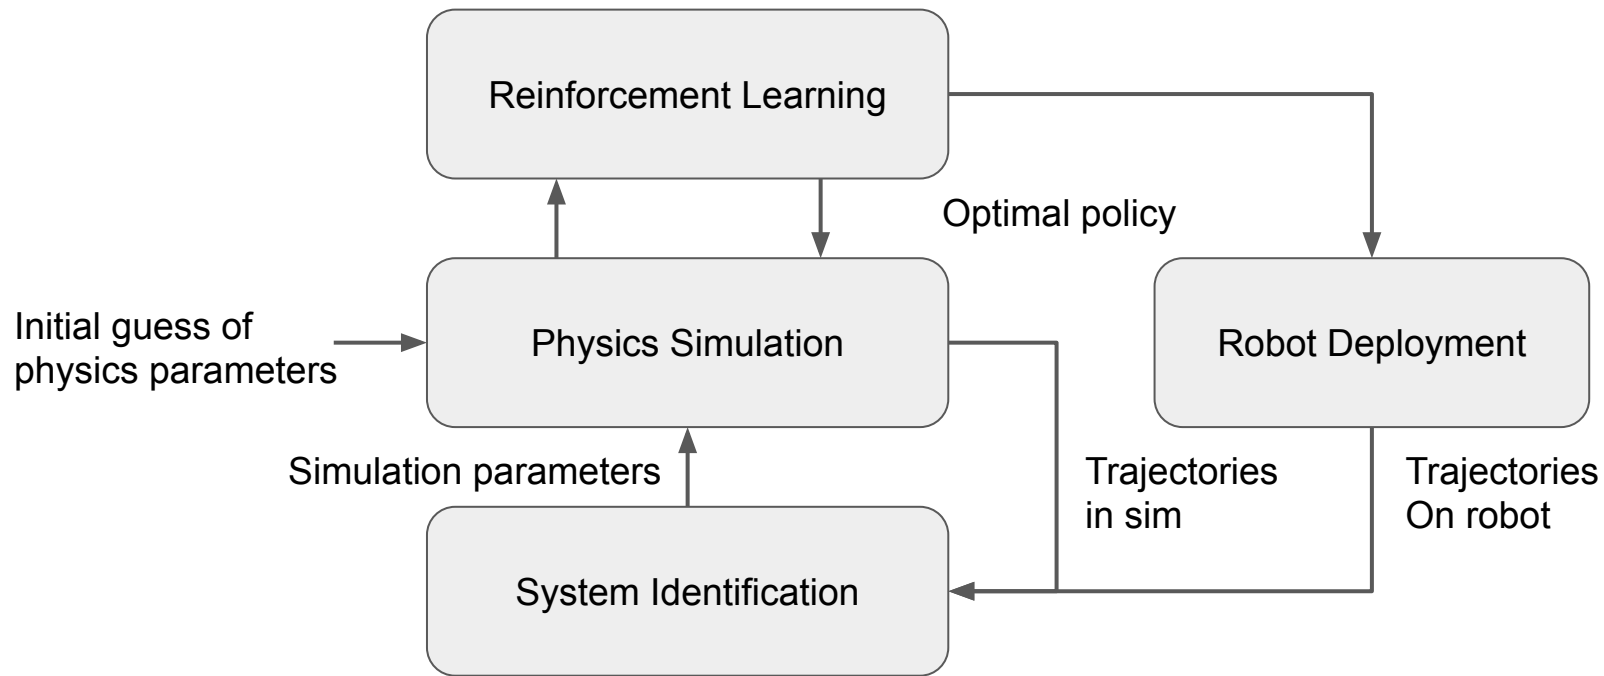


[\[Simulation-based design of dynamic controllers for humanoid balancing, IROS 2016\]](#)

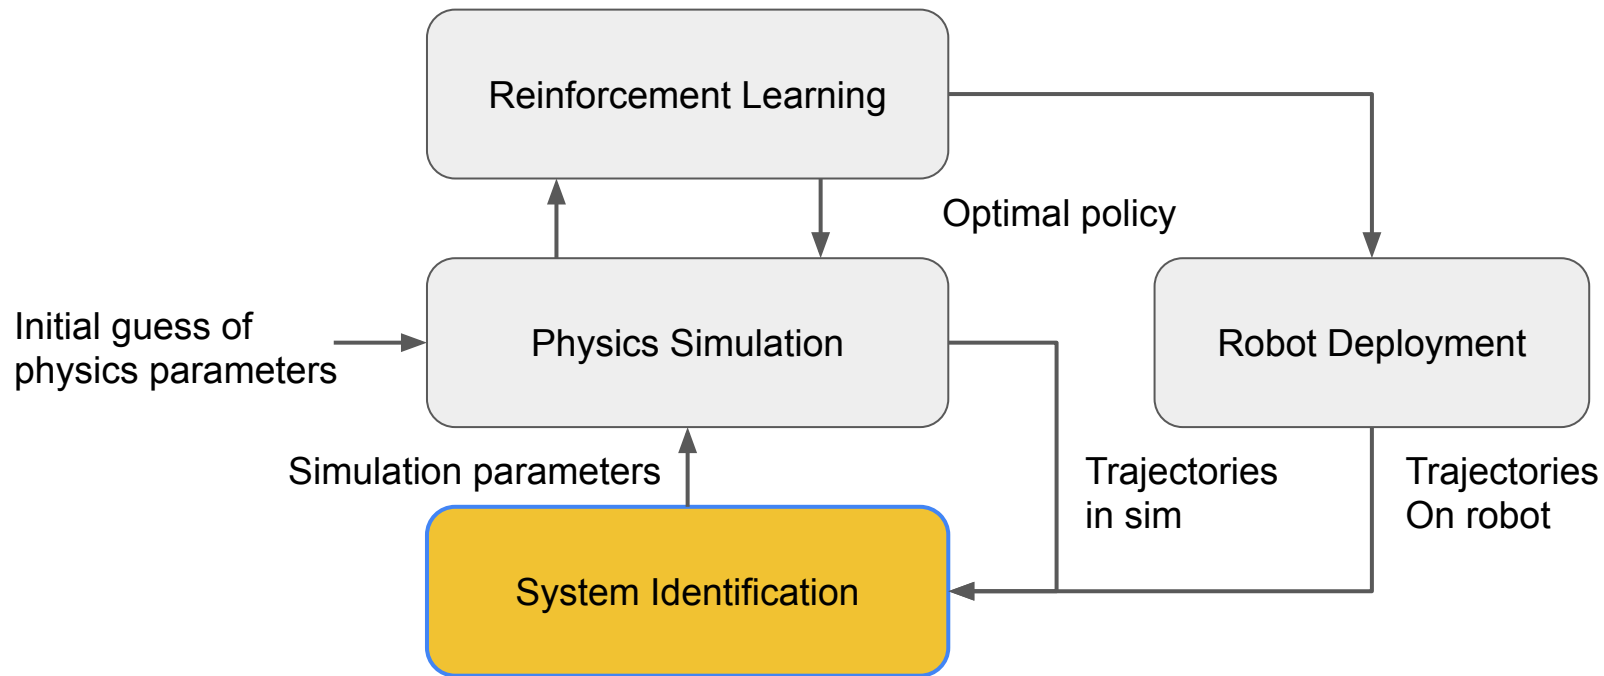


[\[Simulation-based design of dynamic controllers for humanoid balancing, IROS 2016\]](#)





[\[Simulation-based design of dynamic controllers for humanoid balancing, IROS 2016\]](#)



[\[Simulation-based design of dynamic controllers for humanoid balancing, IROS 2016\]](#)


# Automatic System Identification

- Measure sim-to-real discrepancy

$$\theta = \arg \min \frac{1}{n} \sum_{i=1}^n \int_0^{T+1} \|\tilde{\mathbf{q}}_i(t) - \mathbf{q}_i(t; \theta)\|_{\mathbf{W}}^2 dt$$

- Optimize the physics parameters
  - [Covariance Matrix Adaptation-Evolution Strategy](#)

Latency



**Ground truth physical  
parameter:**


Latency = 5ms

Actuator strength = 10nm


Actuator  
strength

**Randomly sampled  
physical parameter:**  
Latency = 1ms  
Actuator strength = 2Nm


**Sim trajectory:**




**Real trajectory:**



**Loss:** 
$$\frac{1}{n} \sum_{i=1}^n \int_0^{T+1} \|\tilde{\mathbf{q}}_i(t) - \mathbf{q}_i(t; \boldsymbol{\theta})\|_{\mathbf{W}}^2 dt$$




Latency




Actuator  
strength

Latency



Actuator  
strength


Latency



Actuator  
strength

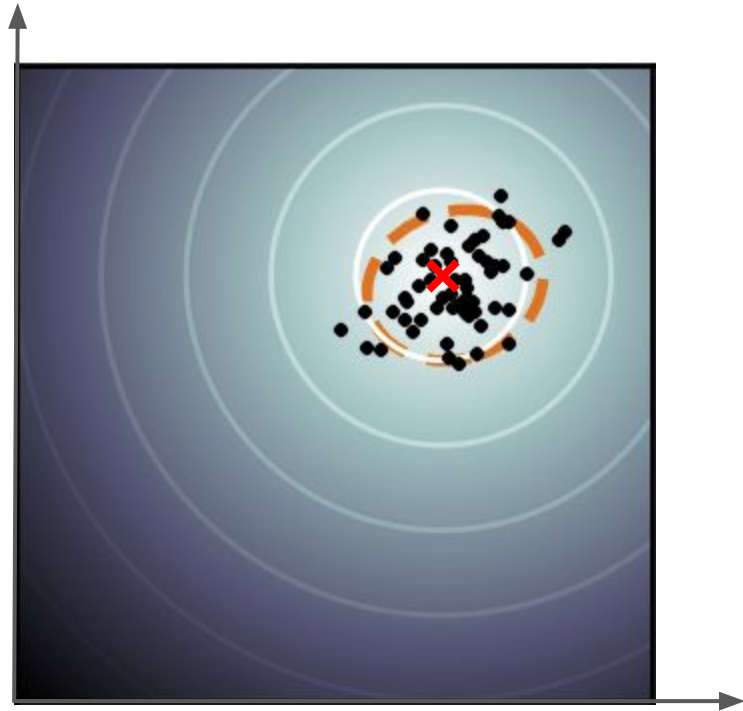


Latency




Actuator  
strength

Latency




Actuator  
strength

Latency

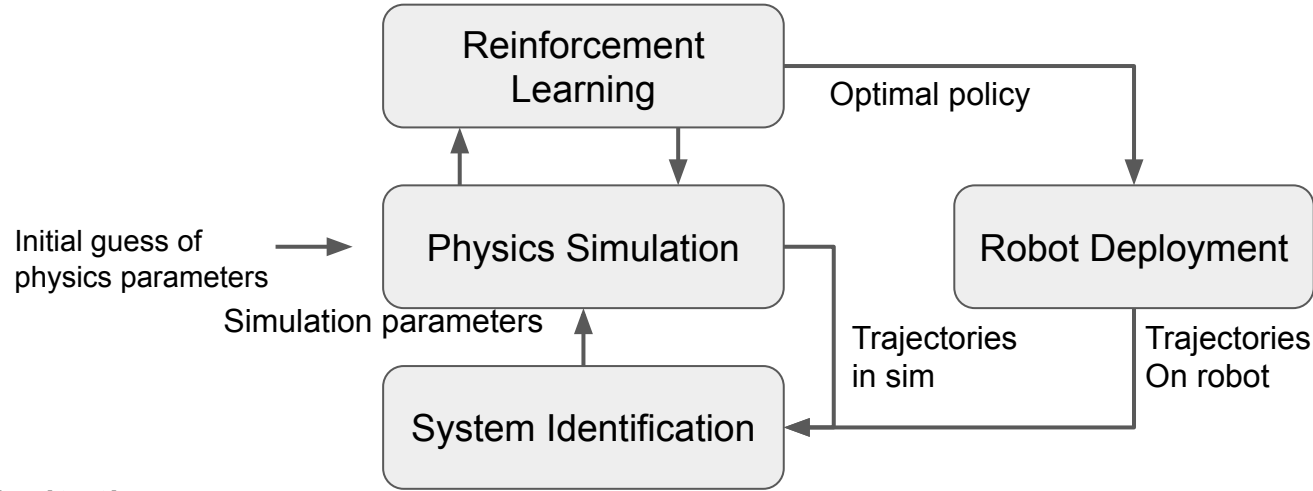


Actuator  
strength



[\[Simulation-based design of dynamic controllers for humanoid balancing, IROS 2016\]](#)

# Automatic System Identification



- Limitations


- Manual selection of physical parameters needed
- Do not work if sim and real trajectory diverge too quickly
- Not account for unmodeled dynamics
- Physical parameters overfit

# Domain Randomization

# Domain Randomization

- Original objective: reward maximization

$$\mathbb{E}_{\tau \sim p(\tau|\pi)} \left[ \sum_{t=0}^{T-1} r(s_t, a_t) \right]$$



# Domain Randomization


- Original objective: reward maximization

$$\mathbb{E}_{\tau \sim p(\tau|\pi)} \left[ \sum_{t=0}^{T-1} r(s_t, a_t) \right]$$

- New objective with domain randomization


$$\mathbb{E}_{\mu \sim \rho_{\mu}} \left[ \mathbb{E}_{\tau \sim p(\tau|\pi, \mu)} \left[ \sum_{t=0}^{T-1} r(s_t, a_t) \right] \right]$$

Physical parameters




Parameter	Range
Link Mass	$[0.25, 4] \times$ default mass of each link
Joint Damping	$[0.2, 20] \times$ default damping of each joint
Puck Mass	$[0.1, 0.4] \text{ kg}$
Puck Friction	$[0.1, 5]$
Puck Damping	$[0.01, 0.2] \text{ N s/m}$
Table Height	$[0.73, 0.77] \text{ m}$
Controller Gains	$[0.5, 2] \times$ default gains
Action Timestep $\lambda$	$[125, 1000] \text{ s}^{-1}$






our method

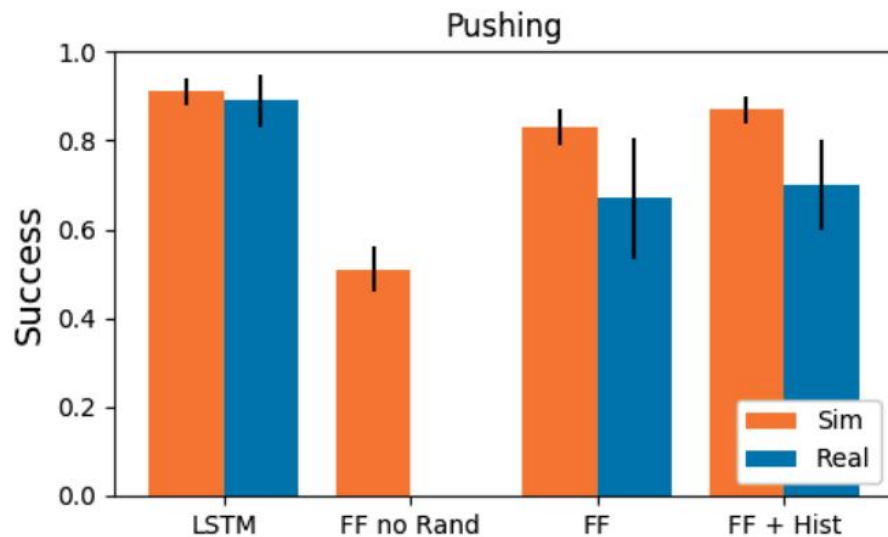



no randomization  
during training


# Memory (LSTM) in sim-to-real




# Memory (LSTM) in sim-to-real



[\[Sim-to-Real Transfer of Robotic Control with Dynamics Randomization, ICRA 2018\]](#)



our method

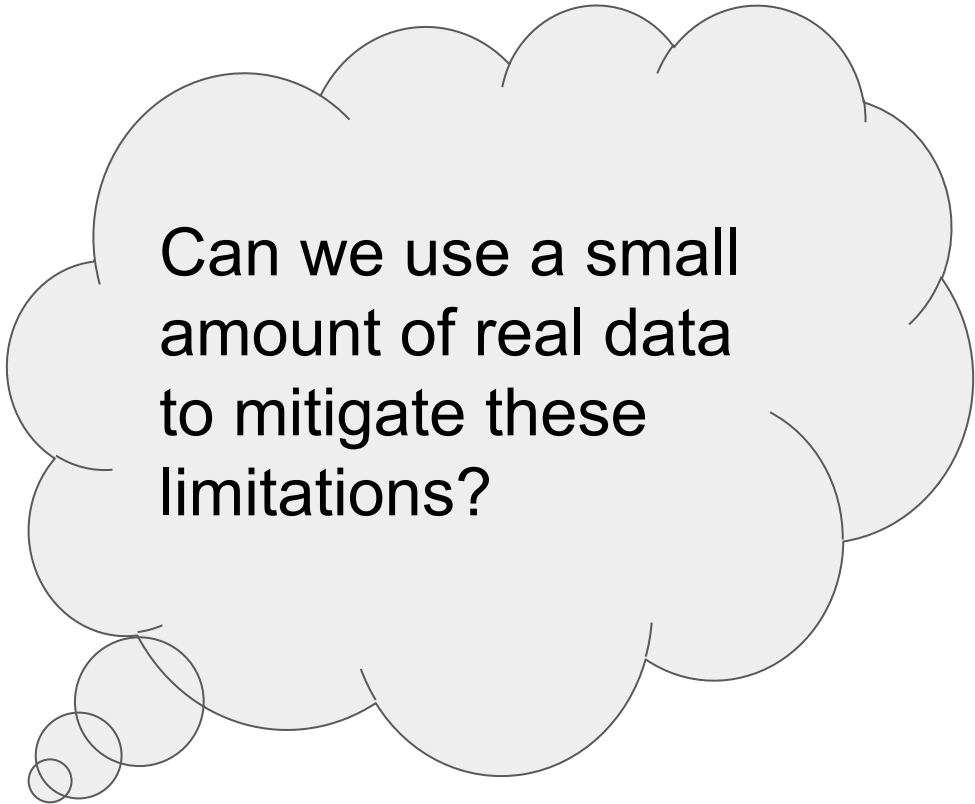


feedforward policy  
(no LSTM)

- Limitations

- Trade optimality for robustness
- Careful tuning needed for the range of randomization

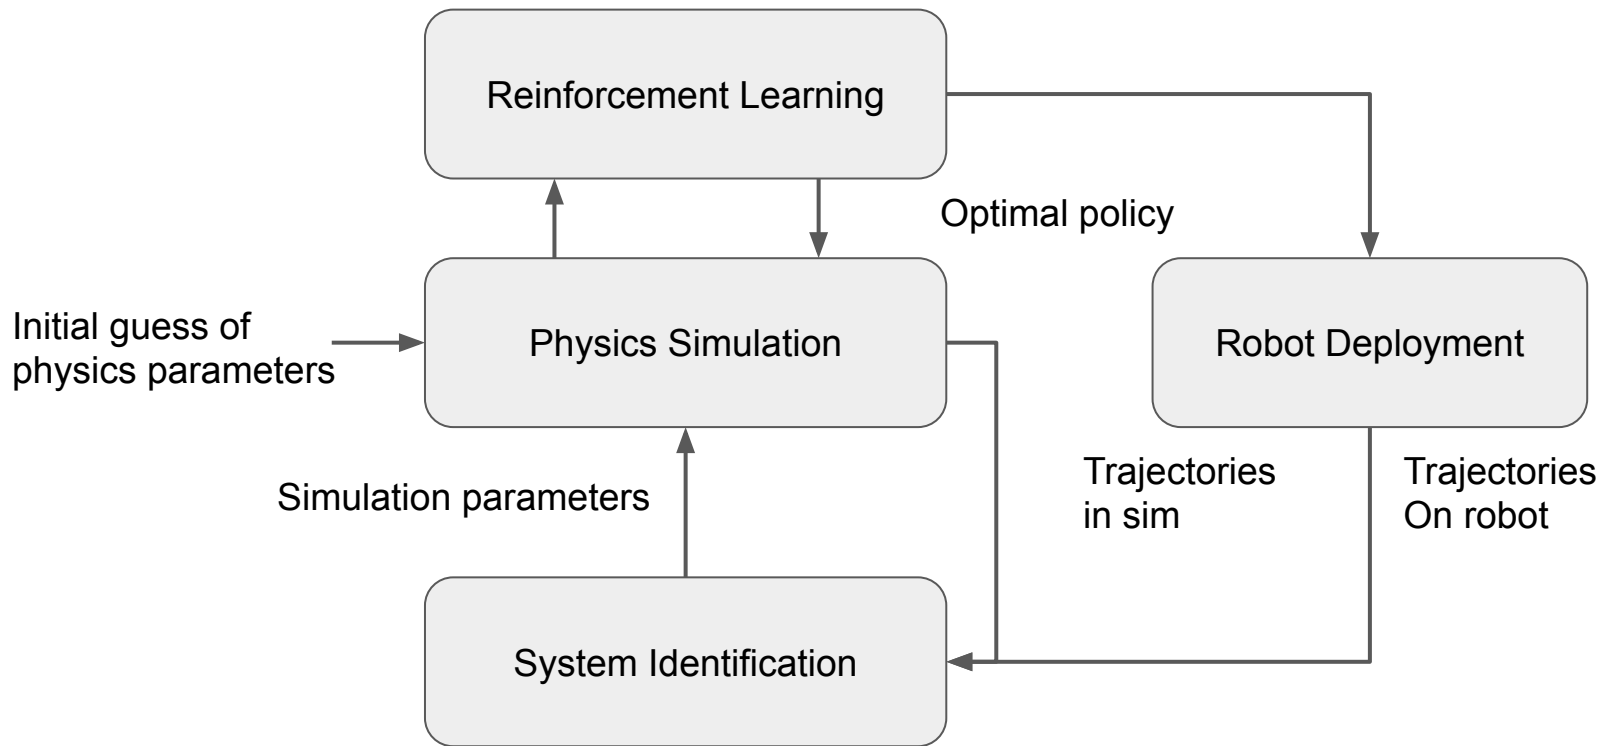
[\[Sim-to-Real Transfer of Robotic Control with Dynamics Randomization, ICRA 2018\]](#)

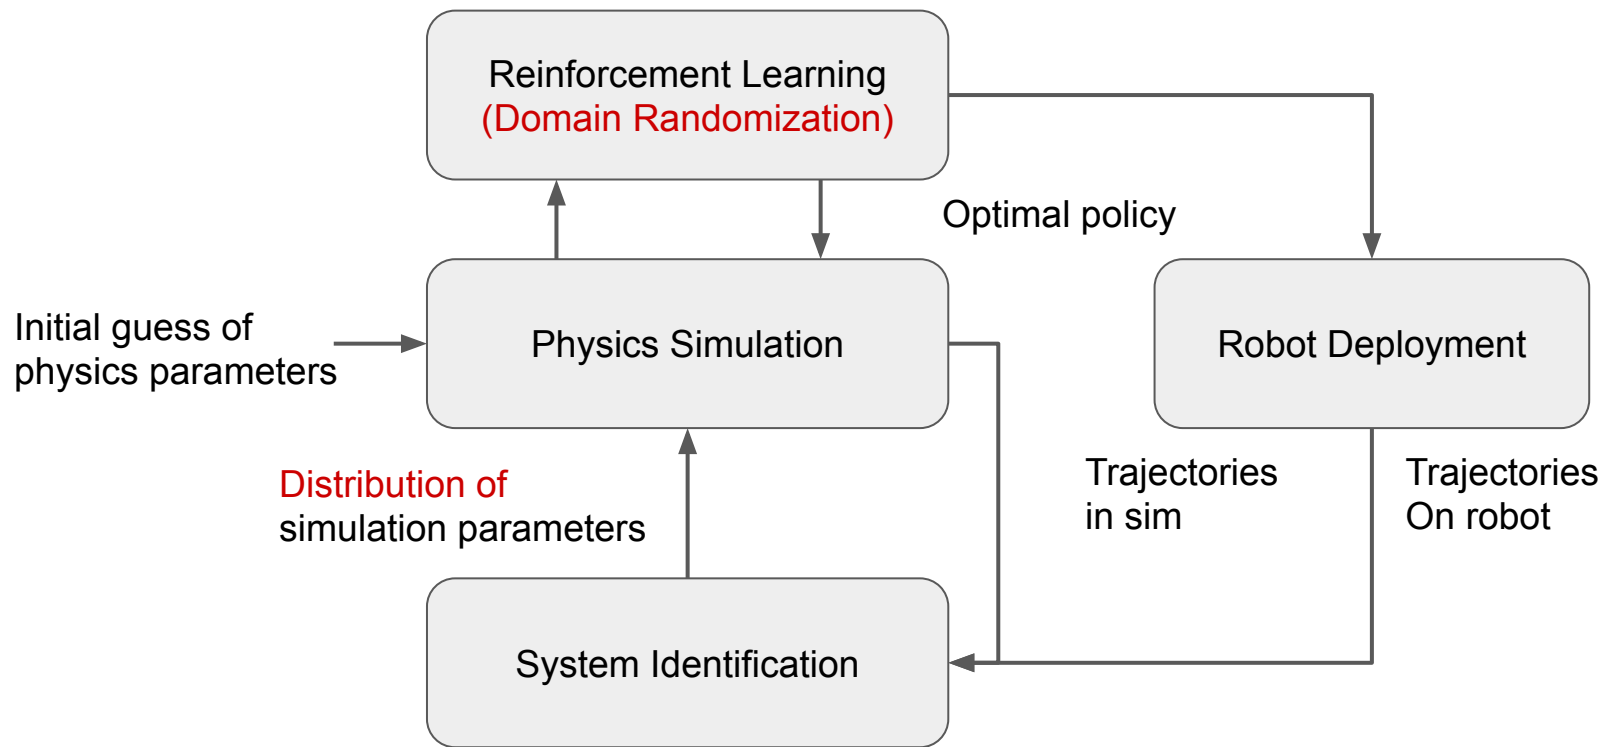


Can we use a small  
amount of real data  
to mitigate these  
limitations?

- Limitations

- Trade optimality for robustness
- Careful tuning needed for the range of randomization








**Swing-peg-in-hole**  
**Simulated environment**

[Closing the Sim-to-Real Loop: Adapting Simulation  
Randomization with Real World Experience, ICRA 2019]



# Domain Adaptation

# Domain Adaptation



[\[Sim-to-Real Transfer for Biped Locomotion, IROS 2019\]](#)

[\[Learning Agile Robotic Locomotion Skills by Imitating Animals, RSS 2020\]](#)

# Domain Adaptation

$$\mu = \begin{bmatrix} \text{Mass} \\ \text{Inertia} \\ \text{Motor Strength} \\ \text{Motor Friction} \\ \text{Latency} \\ \text{Lateral Friction} \\ \text{Etc...} \end{bmatrix}$$

[\[Sim-to-Real Transfer for Biped Locomotion, IROS 2019\]](#)

[\[Learning Agile Robotic Locomotion Skills by Imitating Animals, RSS 2020\]](#)


# Domain Adaptation



[\[Sim-to-Real Transfer for Biped Locomotion, IROS 2019\]](#)

[\[Learning Agile Robotic Locomotion Skills by Imitating Animals, RSS 2020\]](#)


# Domain Adaptation



[\[Sim-to-Real Transfer for Biped Locomotion, IROS 2019\]](#)

[\[Learning Agile Robotic Locomotion Skills by Imitating Animals, RSS 2020\]](#)


# Domain Adaptation



[\[Sim-to-Real Transfer for Biped Locomotion, IROS 2019\]](#)

[\[Learning Agile Robotic Locomotion Skills by Imitating Animals, RSS 2020\]](#)

# Domain Adaptation



$$\mathbf{z}^* = \arg \max_{\mathbf{z}} \mathbb{E}_{\tau \sim p^*(\tau | \pi, \mathbf{z})} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right]$$

[\[Sim-to-Real Transfer for Biped Locomotion, IROS 2019\]](#)

[\[Learning Agile Robotic Locomotion Skills by Imitating Animals, RSS 2020\]](#)

# Domain Adaptation vs. Domain Randomization

Dog Pace



No Randomization



Randomization



Domain Adaptation (Ours)



# Domain Adaptation vs. Domain Randomization

Dog Spin



No Randomization



Randomization



Domain Adaptation (Ours)

- Limitations

- Policy is not updated
- The latent space may not contain the optimal vector for the real world
- Performance not necessarily improve with more real data

# Meta Learning

# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.

# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.


- Real robot: mass = unknown

●  
 $\theta_{\text{rea}}$   
|

# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.


- Real robot: mass = unknown
- Sim task 1: mass = 7.0kg
- Sim task 2: mass = 9.2kg
- Sim task 3: mass = 9.8kg



# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.


- Real robot: mass = unknown
- Sim task 1: mass = 7.0kg
- Sim task 2: mass = 9.2kg
- Sim task 3: mass = 9.8kg



# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.


- Real robot: mass = unknown
- Sim task 1: mass = 7.0kg
- Sim task 2: mass = 9.2kg
- Sim task 3: mass = 9.8kg



# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.

- Real robot: mass = unknown
- Sim task 1: mass = 7.0kg
- Sim task 2: mass = 9.2kg
- Sim task 3: mass = 9.8kg






# MAML for Sim-to-Real Gap

Learn mostly in simulation; quickly adapt to the real world with few real rollouts.

- Real robot: mass = unknown
- Sim task 1: mass = 7.0kg
- Sim task 2: mass = 9.2kg
- Sim task 3: mass = 9.8kg



# MAML for Sim-to-Real Gap

- Fine-tuning

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$$

- Meta-training

$$\min_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

---

**Algorithm 1** Model-Agnostic Meta-Learning

---

**Require:**  $p(\mathcal{T})$ : distribution over tasks

**Require:**  $\alpha, \beta$ : step size hyperparameters

- 1: randomly initialize  $\theta$
  - 2: **while** not done **do**
  - 3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$
  - 4:   **for all**  $\mathcal{T}_i$  **do**
  - 5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples
  - 6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
  - 7:   **end for**
  - 8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
  - 9: **end while**
-

# MAML for Sim-to-Real Gap

- Fine-tuning

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$$

- Meta-training

$$\min_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

---

**Algorithm 1** Model-Agnostic Meta-Learning

---

**Require:**  $p(\mathcal{T})$ : distribution over tasks

**Require:**  $\alpha, \beta$ : step size hyperparameters

1: randomly initialize  $\theta$

2: **while** not done **do**

3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$

4:   **for all**  $\mathcal{T}_i$  **do**

5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples

6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$

7:   **end for**

8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$

9: **end while**

---

# MAML for Sim-to-Real Gap

- Fine-tuning

$$\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$$

- Meta-training

$$\min_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$$

---

---

## Algorithm 1 Model-Agnostic Meta-Learning

---

**Require:**  $p(\mathcal{T})$ : distribution over tasks

**Require:**  $\alpha, \beta$ : step size hyperparameters

1: randomly initialize  $\theta$

2: **while** not done **do**

3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$

4:   **for all**  $\mathcal{T}_i$  **do**

5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples

6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$

7:   **end for**

8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$

9: **end while**

---

# Gradient Estimation using Evolution Strategy (ES)

$$\nabla_{\theta} \tilde{f}_{\sigma}(\theta) = \frac{1}{\sigma} \mathbb{E}_{\mathbf{g} \sim \mathcal{N}(0, \mathbf{I}_d)} [\underbrace{f(\theta + \sigma \mathbf{g})}_{\text{Return of perturbed policy}} \underbrace{\mathbf{g}}_{\text{Perturbation direction}}]$$

# Gradient Estimation using Evolution Strategy (ES)

$$\nabla_{\theta} \tilde{f}_{\sigma}(\theta) = \frac{1}{\sigma} \mathbb{E}_{\mathbf{g} \sim \mathcal{N}(0, \mathbf{I}_d)} [f(\theta + \sigma \mathbf{g}) \mathbf{g}]$$

Estimated policy gradient

# ES-MAML for Reality Gap

---

**Algorithm 1** Model-Agnostic Meta-Learning

---

**Require:**  $p(\mathcal{T})$ : distribution over tasks

**Require:**  $\alpha, \beta$ : step size hyperparameters

1: randomly initialize  $\theta$

2: **while** not done **do**

3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$

4:   **for all**  $\mathcal{T}_i$  **do**

5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples

6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$

7:   **end for**

8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$

9: **end while**

---

**Data:** initial policy  $\theta_0$ , adaptation step size  $\alpha$ ,  
meta step size  $\beta$ , number of queries  $K$

1 **for**  $t = 0, 1, \dots$  **do**

2   Sample  $n$  tasks  $T_1, \dots, T_n$  and iid vectors  
    $\mathbf{g}_1, \dots, \mathbf{g}_n \sim \mathcal{N}(0, \mathbf{I})$ ;

3   **foreach**  $(T_i, \mathbf{g}_i)$  **do**

4      $\mathbf{d}^{(i)} \leftarrow \text{ESGRAD}(f^{T_i}, \theta_t + \sigma \mathbf{g}_i, K, \sigma)$ ;

5      $\theta_t^{(i)} \leftarrow \theta_t + \sigma \mathbf{g}_i + \alpha \mathbf{d}^{(i)}$ ;

6      $v_i \leftarrow f^{T_i}(\theta_t^{(i)})$ ;

7   **end**

8    $\theta_{t+1} \leftarrow \theta_t + \frac{\beta}{\sigma n} \sum_{i=1}^n v_i \mathbf{g}_i$ ;

9 **end**

# ES-MAML for Reality Gap

---

**Algorithm 1** Model-Agnostic Meta-Learning

---

**Require:**  $p(\mathcal{T})$ : distribution over tasks

**Require:**  $\alpha, \beta$ : step size hyperparameters

```
1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  with respect to  $K$  examples
6:     Compute adapted parameters with gradient descent:  $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ 
7:   end for
8:   Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$ 
9: end while
```

---

**Data:** initial policy  $\theta_0$ , adaptation step size  $\alpha$ , meta step size  $\beta$ , number of queries  $K$

```
1 for  $t = 0, 1, \dots$  do
2   Sample  $n$  tasks  $T_1, \dots, T_n$  and iid vectors  $\mathbf{g}_1, \dots, \mathbf{g}_n \sim \mathcal{N}(0, \mathbf{I})$ ;
3   foreach  $(T_i, \mathbf{g}_i)$  do
4      $\mathbf{d}^{(i)} \leftarrow \text{ESGRAD}(f^{T_i}, \theta_t + \sigma \mathbf{g}_i, K, \sigma)$ ;
5      $\theta_t^{(i)} \leftarrow \theta_t + \sigma \mathbf{g}_i + \alpha \mathbf{d}^{(i)}$ ;
6      $v_i \leftarrow f^{T_i}(\theta_t^{(i)})$ ;
7   end
8    $\theta_{t+1} \leftarrow \theta_t + \frac{\beta}{\sigma n} \sum_{i=1}^n v_i \mathbf{g}_i$ ;
9 end
```



Initial Policy

After 30 Episodes

After 50 Episodes



The initial policy shifts to the right.



Domain  
Randomization



PG-MAML




Our Method


# Summary

- Improve simulation
  - System identification
    - Always should be the first step for sim-to-real
    - Identified parameters can be reused
    - Require diverse data and careful experiment design
    - Can be tedious
- Improve policy
  - Domain randomization
    - Simple
    - Zero-shot transfer is often possible
    - Trade-off between robustness and optimality
  - Domain adaptation
    - Small amount of real-world data needed
    - Good experience so far in multiple locomotion projects
  - Meta learning
    - Novel and popular research direction

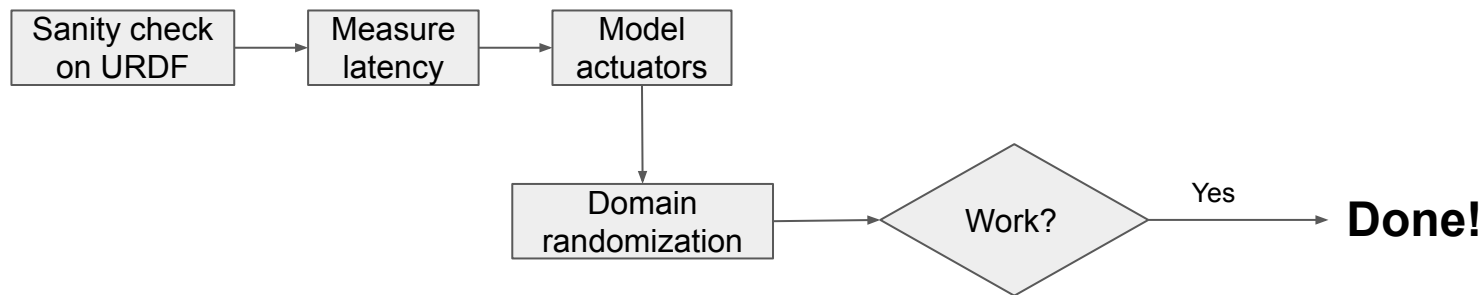
# Pipeline for Sim-to-Real



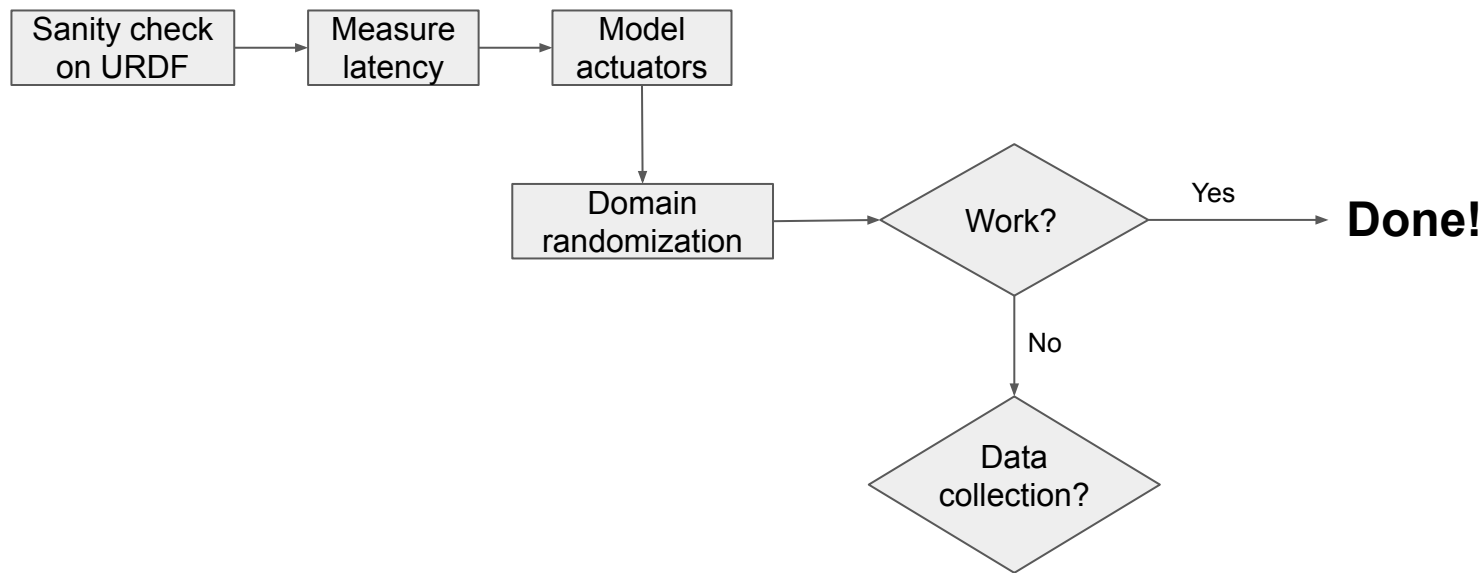
# Pipeline for Sim-to-Real




# Pipeline for Sim-to-Real



# Pipeline for Sim-to-Real

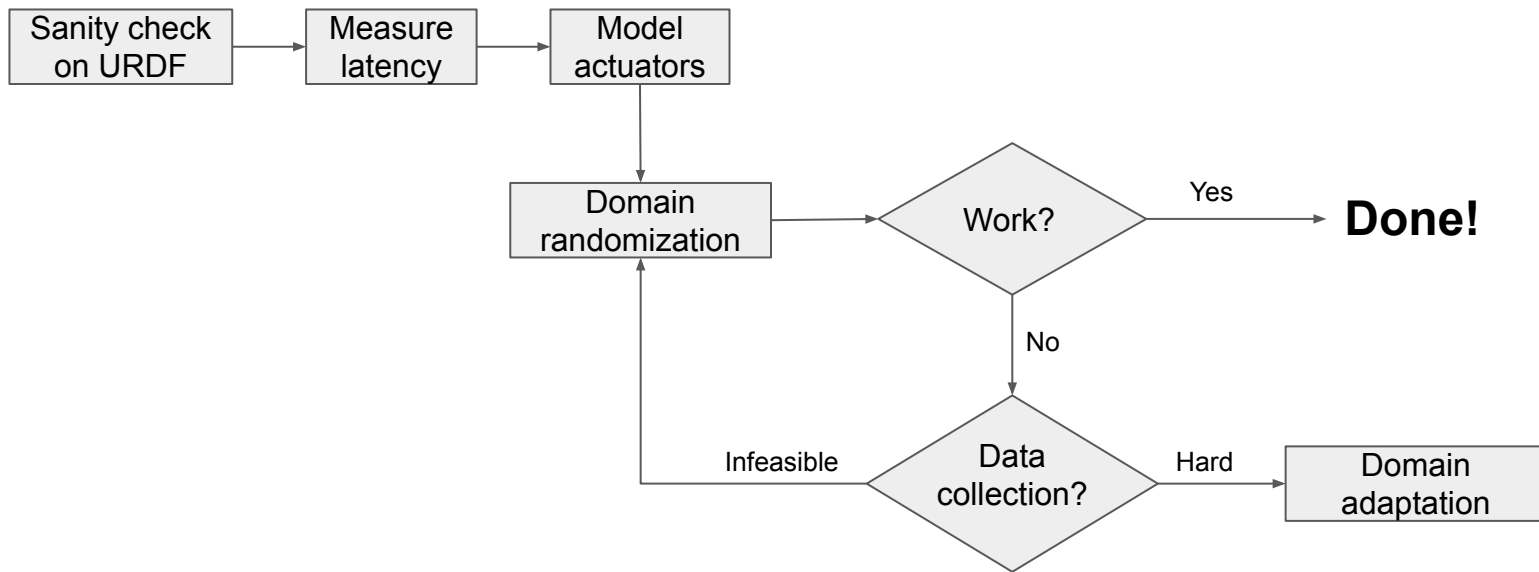


# Pipeline for Sim-to-Real





# Pipeline for Sim-to-Real



# Pipeline for Sim-to-Real

